

Space Weather Forecasting — AI Solar Storm (M-Class Flare) Predictor

Intel AI Global Impact Festival — Proof of Concept (PoC)

Goal: Predict solar storms (X-ray flux > $1e-5$ W/m², M-Class flares) to help protect critical Earth infrastructure (Power Grids, GPS, Aviation).

SDG Alignment:

-  **UN SDG 9** — Industry, Innovation & Infrastructure
-  **UN SDG 13** — Climate Action

Environment note: Built for Python 3.13. This notebook intentionally uses **standard scikit-learn** (NOT `scikit-learn-intelex`), since `scikit-learn-intelex` currently fails to install / is incompatible with Python 3.13 environments.

Data policy: This notebook uses **100% real data** from the NOAA SWPC GOES satellite API. There is **no simulated/synthetic data** anywhere in this pipeline. If the NOAA API is unreachable, the notebook raises a clear error instead of inventing data — for a space-weather safety application, silently substituting fake data for a failed real-time feed would be misleading.

0. Imports & Setup

```
import sys
import warnings
from datetime import datetime, timedelta

import numpy as np
import pandas as pd
import requests
import matplotlib.pyplot as plt

from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report

warnings.filterwarnings("ignore")
```

1. Data Sourcing (100% Real NOAA Data Only)

Note on NOAA API change: NOAA SWPC previously served a 30-day X-ray file, but their public "primary" JSON directory no longer offers one — the longest real-time X-ray window they publish is **7 days**:

```
https://services.swpc.noaa.gov/json/goes/primary/xrays-7-day.json
```

This notebook uses **only** that real NOAA data. If the API is unreachable (network down, endpoint renamed again, rate-limited, etc.), the next cell raises a clear `RuntimeError` and stops rather than inventing data.

```
NOAA_URL = "https://services.swpc.noaa.gov/json/goes/primary/xrays-7-day.json"
STORM_THRESHOLD = 1e-5 # M-Class flare threshold (W/m^2), per NOAA scale

def fetch_real_noaa_data(url: str) -> pd.DataFrame:
    """Fetch real-time X-ray flux data from NOAA SWPC (max real window: 7
    days)."""
    response = requests.get(url, timeout=10, headers={"User-Agent":
    "Mozilla/5.0"})
    response.raise_for_status()
    raw = response.json()

    df = pd.DataFrame(raw)

    # NOAA's JSON uses 'time_tag' and 'flux' (with an 'energy' column to
    # distinguish 0.05-0.4nm vs 0.1-0.8nm channels). We want the long
    channel
    # (0.1-0.8nm), which is the standard GOES X-ray Class channel.
    if "energy" in df.columns:
        df = df[df["energy"] == "0.1-0.8nm"]

    df = df.rename(columns={"time_tag": "time", "flux": "long_flux"})
    df["time"] = pd.to_datetime(df["time"])
    df = df[["time", "long_flux"]].dropna()
    df =
df.sort_values("time").drop_duplicates(subset="time").reset_index(drop=True)

    if df.empty:
        raise ValueError("NOAA response parsed but contained no usable
        rows.")

    return df
```

```

print("=" * 80)
print("STEP 1: Fetching Solar X-ray Flux Data (NOAA SWPC GOES
Satellite)")
print("=" * 80)

try:
    df = fetch_real_noaa_data(NOAA_URL)
    span_days = (df["time"].max() - df["time"].min()).total_seconds() /
86400
    data_source = f"REAL-TIME NOAA SWPC API (xrays-7-day.json) -
{len(df)} records"
    print(f"[OK] Successfully fetched {len(df)} REAL records from NOAA
SWPC.")
    print(f"[INFO] Real data span: {span_days:.1f} days "
          f"(NOAA's public API no longer offers a 30-day X-ray window; "
          f"7 days is the longest real-time feed currently available).")
except Exception as e:
    print(f"[ERROR] NOAA API fetch failed: {type(e).__name__}: {e}")
    print("[ERROR] This notebook uses ONLY real NOAA data with no
simulated "
          "fallback, per project requirements. Cannot continue without a
"
          "live data source. Please check your internet connection / NOAA
"
          "SWPC service status (https://www.swpc.noaa.gov) and try
again.")
    raise RuntimeError("Aborting: no real NOAA data available.") from e

print(f"[INFO] Data source used: {data_source}")
print(f"[INFO] Records: {len(df)} | Range: {df['time'].min()} ->
{df['time'].max()}")
df.head()

```

```

=====
STEP 1: Fetching Solar X-ray Flux Data (NOAA SWPC GOES Satellite)
=====

```

```

[OK] Successfully fetched 10077 REAL records from NOAA SWPC.
[INFO] Real data span: 7.0 days (NOAA's public API no longer offers a 30-day X-
ray window; 7 days is the longest real-time feed currently available).
[INFO] Data source used: REAL-TIME NOAA SWPC API (xrays-7-day.json) - 10077
records
[INFO] Records: 10077 | Range: 2026-08-18 14:07:00+00:00 -> 2026-08-25
14:03:00+00:00

```

	time	long_flux
0	2026-08-18 14:07:00+00:00	0.000002
1	2026-08-18 14:08:00+00:00	0.000002
2	2026-08-18 14:09:00+00:00	0.000002
3	2026-08-18 14:10:00+00:00	0.000002
4	2026-08-18 14:11:00+00:00	0.000002

2. Feature Engineering — TRUE Forecasting (Not Nowcasting)

⚠ Critical fix: NOAA's `xrays-7-day.json` is published at **~1-minute cadence**, not hourly. A naive `shift(1) - shift(5)` therefore lags by 1–5 *minutes*, and X-ray flux barely changes in under a minute — so a model trained that way isn't forecasting, it's just echoing the almost-identical previous reading. That's why an earlier version of this notebook produced a suspicious ~99.85% accuracy: the "prediction" task was trivial.

This version fixes it properly:

1. Detects the *actual* sampling interval from the real NOAA timestamps (no hardcoded assumption).
2. Builds lag features from genuinely-past readings.
3. Sets the **target ~60 minutes into the future** (a real forecast horizon), instead of the current/near-current value — so the model must genuinely predict, not copy.

```
print("=" * 80)
print("STEP 2: Feature Engineering (Lag Features + TRUE Forecast Target)")
print("=" * 80)

df = df.sort_values("time").reset_index(drop=True)

# --- Detect the REAL sampling interval of the NOAA data -----
# -----
time_diffs = df["time"].diff().dt.total_seconds().dropna()
dt_seconds = time_diffs.median()
dt_minutes = dt_seconds / 60
print(f"[INFO] Detected NOAA sampling interval: ~{dt_minutes:.1f} minute(s) per reading.")

# Forecast horizon: predict ~60 minutes ahead (genuine early-warning lead time)
```

```

forecast_horizon_steps = max(1, round(60 / dt_minutes))
print(f"[INFO] Forecast horizon: {forecast_horizon_steps} step(s) "
      f"(~{forecast_horizon_steps * dt_minutes:.0f} minutes ahead of the "
      f"most recent reading used as input).")

# --- Lag features: genuinely PAST readings -----
-----
for lag in range(1, 6):
    df[f"flux_lag_{lag}"] = df["long_flux"].shift(lag)

# --- Target: will an M-Class storm be happening ~60 minutes from NOW? ---
-----
# We look FORWARD (shift with a negative value) so the target represents
the
# future, not the present – this is what makes it a forecast.
df["is_storm_now"] = (df["long_flux"] > STORM_THRESHOLD).astype(int)
df["is_storm"] = df["is_storm_now"].shift(-forecast_horizon_steps)

model_df = df.dropna().reset_index(drop=True)
model_df["is_storm"] = model_df["is_storm"].astype(int)

feature_cols = [f"flux_lag_{lag}" for lag in range(1, 6)]
X = model_df[feature_cols]
y = model_df["is_storm"]

print(f"[INFO] Feature columns (past readings): {feature_cols}")
print(f"[INFO] Target: is_storm ~{forecast_horizon_steps *
dt_minutes:.0f} minutes into the future")
print(f"[INFO] Storm class balance -> Storm=1: {y.sum()} | Quiet=0: {(y
== 0).sum()}")

if y.nunique() < 2:
    print("[ERROR] The real NOAA data currently fetched contains only ONE "
          "class (no M-class storms occurred/forecast in this window, "
          "i.e. "
          "the sun has been quiet). A binary classifier needs both "
          "classes "
          "to train. Since this notebook uses real data only (no "
          "synthetic "
          "samples), it cannot proceed right now.")
    print("[INFO] This is expected/normal behavior during quiet solar "
          "periods "
          "- try again during a period with more solar activity, or check "
          "https://www.swpc.noaa.gov for recent flare history.")
    raise RuntimeError("Aborting: only one class present in target (no "
                       "storms in window).")

model_df.tail()

```



```
y_pred = model.predict(X_test)
y_proba = model.predict_proba(X_test)[:, 1]
print("[OK] Model trained.")
```

```
=====
STEP 3: Training RandomForestClassifier (80/20 split)
=====
```

```
[OK] Model trained.
```

💡 Intel Tie-In (for pitch deck)

This `RandomForestClassifier` is our PoC baseline trained with standard scikit-learn. For the production / real-time deployment story, this is where **Intel technology** becomes the key differentiator:

1. Intel OpenVINO Toolkit

- Convert the trained scikit-learn `RandomForest` into an OpenVINO IR (Intermediate Representation) using OpenVINO's model conversion tools (or export via ONNX → OpenVINO Model Optimizer).
- Quantize the model (INT8) with OpenVINO's Neural Network Compression Framework (NNCF) to shrink model size and boost inference speed with minimal accuracy loss.

2. Edge Deployment at Ground Stations

- Ground-based magnetometer / X-ray monitoring stations often run on low-power industrial edge hardware (Intel Core / Atom CPUs), not GPU servers. OpenVINO's CPU-optimized inference engine lets this forecasting model run in near real-time directly on that hardware, without needing constant cloud connectivity — critical when a geomagnetic storm may itself disrupt satellite comms links.

3. Why this matters for SDG 9 & 13

- Faster, on-site inference = faster storm warnings pushed to power grid operators, aviation route planners, and GPS-dependent systems **before** the storm's effects arrive (X-rays reach Earth in ~8 minutes, but grid/GPS operators need lead time to act).
- Low-power edge inference also reduces the energy/data-transfer footprint of the early-warning network itself — supporting sustainable, resilient infrastructure (SDG 9) and climate/disaster resilience (SDG 13).

4. Model Evaluation

⚠️ **Note on metrics:** M-Class storms are rare events, so this is a heavily imbalanced

classification problem. **Accuracy alone can be misleading** here — a model that always predicts "quiet" could still score >95% accuracy while being useless. We report Confusion Matrix and ROC-AUC alongside the Classification Report for an honest picture.

```
from sklearn.metrics import confusion_matrix, roc_auc_score,
RocCurveDisplay

print("=" * 80)
print("STEP 4: Model Evaluation")
print("=" * 80)

acc = accuracy_score(y_test, y_pred)
print(f"\nAccuracy Score: {acc:.4f} (see note above - can be misleading
on imbalanced data)\n")
print("Classification Report:")
print(classification_report(y_test, y_pred, target_names=["Quiet (0)",
"Storm (1)"], zero_division=0))

cm = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
print(f"
      Predicted Quiet    Predicted Storm")
print(f"Actual Quiet      {cm[0][0]:>15}    {cm[0][1]:>15}")
print(f"Actual Storm      {cm[1][0]:>15}    {cm[1][1]:>15}")

if y_test.nunique() > 1:
    auc = roc_auc_score(y_test, y_proba)
    print(f"\nROC-AUC Score: {auc:.4f} (0.5 = random guessing, 1.0 =
perfect)")
    RocCurveDisplay.from_predictions(y_test, y_proba)
    plt.title("ROC Curve - Storm Forecast (~60 min ahead)")
    plt.tight_layout()
    plt.savefig("roc_curve.png", dpi=150)
    plt.show()
    print("[OK] ROC curve saved as 'roc_curve.png'")

feat_importance = pd.Series(model.feature_importances_,
index=feature_cols).sort_values(ascending=False)
print("\nFeature Importances (which lag matters most for forecasting):")
print(feat_importance.to_string())
```

```
=====
STEP 4: Model Evaluation
=====
```

```
Accuracy Score: 0.7509 (see note above - can be misleading on imbalanced data)
```

```
Classification Report:
```


```
precision    recall  f1-score   support
```


Quiet (0)	0.98	0.76	0.86	1940
Storm (1)	0.06	0.49	0.11	63
accuracy			0.75	2003
macro avg	0.52	0.63	0.48	2003
weighted avg	0.95	0.75	0.83	2003

Confusion Matrix:

	Predicted Quiet	Predicted Storm
Actual Quiet	1473	467
Actual Storm	32	31

ROC-AUC Score: 0.7615 (0.5 = random guessing, 1.0 = perfect)



[OK] ROC curve saved as 'roc_curve.png'

Feature Importances (which lag matters most for forecasting):

flux_lag_5	0.253717
flux_lag_1	0.223717
flux_lag_2	0.191670
flux_lag_3	0.183519
flux_lag_4	0.147377

5. Visualization

Plot X-ray flux over time with the M-Class storm threshold marked. Saved as solar_storm_graph.png .

```
print("=" * 80)
print("STEP 5: Generating Visualization -> solar_storm_graph.png")
print("=" * 80)

fig, ax = plt.subplots(figsize=(14, 6))

ax.plot(df["time"], df["long_flux"], color="#1f77b4", linewidth=1.2,
        label="Real GOES Long-Wavelength X-ray Flux (0.1-0.8nm), NOAA SWPC")

ax.axhline(y=STORM_THRESHOLD, color="red", linestyle="--", linewidth=1.5,
           label="M-Class Storm Threshold (1e-5 W/m\u00b2)")


storms = df[df["long_flux"] > STORM_THRESHOLD]
if not storms.empty:
    ax.scatter(storms["time"], storms["long_flux"], color="red", s=25,
               zorder=5, label="Observed Storm Events")

ax.set_yscale("log")
ax.set_xlabel("Time (UTC)")
ax.set_ylabel("X-ray Flux (W/m\u00b2) [log scale]")
ax.set_title(f"Space Weather Forecasting: Solar X-ray Flux\n({span_days:.1f}-Day Real NOAA Window)\nData Source: {data_source} | Forecast horizon: ~{forecast_horizon_steps * dt_minutes:.0f} min ahead")
ax.legend(loc="upper right")
ax.grid(True, which="both", linestyle=":", alpha=0.5)
fig.autofmt_xdate()

plt.tight_layout()
plt.savefig("solar_storm_graph.png", dpi=150)
print("[OK] Plot saved as 'solar_storm_graph.png'")
plt.show()

print("\n" + "=" * 80)
print("PoC COMPLETE. Ready for Intel AI Global Impact Festival pitch.")
print("=" * 80)

=====
STEP 5: Generating Visualization -> solar_storm_graph.png
=====
[OK] Plot saved as 'solar_storm_graph.png'
```



=====
PoC COMPLETE. Ready for Intel AI Global Impact Festival pitch.
=====

Honest Limitations (for pitch / judges Q&A)

- **Data window is short (real NOAA data only spans ~7 days)** — not enough history to capture rare M-class events reliably. More historical real data (e.g. via NOAA's SWPC data archive / NCEI, rather than the live 7-day feed) would substantially improve training.
- **Forecasting solar flares 60 minutes ahead from X-ray flux history alone is a genuinely hard, actively-researched problem** — flares can onset abruptly. Don't be surprised if storm-class recall is modest; that's expected and honest, not a bug.
- **Class imbalance** — M-class storms are rare relative to quiet periods, which is why Accuracy alone is not a reliable metric here (see ROC-AUC and Confusion Matrix above instead).
- **Next steps for a stronger model:** incorporate additional real NOAA channels (e.g. proton flux, magnetometer data), use a longer historical archive, and/or engineer trend/rate-of-change features (not just raw lag values).